Confluent Platform Quick Start (Docker)

This quick start shows you how to get up and running with Confluent Platform and its main components using Docker containers. This quick start demonstrates both the basic and most powerful capabilities of Confluent Platform, including using Control Center for topic management and event stream processing using KSQL. In this quick start, you create Apache Kafka® topics, use Kafka Connect to generate mock data to those topics, and create KSQL streaming queries on those topics. You then go to Control Center to monitor and analyze the event streaming queries.

See also

You can also run an automated version of this quick start (https://github.com/confluentinc/examples/tree/5.4.0-post/cp-quickstart/README.md) designed for Confluent Platform local installs.

Prerequisites:

- Docker:

Issue: Cannot

- Docker version 1.11 or later is installed and running (https://docs.docker.com/engine/installation/)

the default Docker memory allocation is 2 GB. You can change the default allocation to 8 GB in **Docker** > **Preferences** > **Advanced**.

- Git (https://git-scm.com/downloads).
- Internet connectivity.
- Ensure you are on an Operating System (../installation/versions-interoperability.html#operating-systems) currently supported by Confluent Platform.
- Networking and Kafka on Docker: Configure your hosts and ports to allow both internal and external components to the Docker network to communicate. For more details, see this article (https://rmoff.net/2018/08/02/kafka-listeners-explained/).

Step 1: Download and Start Confluent Platform Using Docker

1. Clone the Confluent Platform Docker Images GitHub Repository and check out the `5.4.0-post` branch.

```
git clone https://github.com/confluentinc/examples
cd examples
git checkout 5.4.0-post
```


Issue: Cannot

2. Navigate to `cp-all-in-one` examples directory.

```
(/current/tutorials/cp-all-in-one/docker-docker-docker-quickstart.html)
```

3. Start Confluent Platform specifying two options: (`-d`) to run in detached mode and (`--build`) to build the Kafka Connect image with the source connector `kafka-connect-datagen` from Confluent Hub (https://www.confluent.io/connector/kafka-connect-datagen/).

Important

You must allocate a minimum of 8 GB of Docker memory resource. The default memory allocation on Docker Desktop for Mac is 2 GB and must be changed.

```
docker-compose up -d --build
```

This starts Confluent Platform with separate containers for all Confluent Platform components. Your output should resemble the following:

```
Creating network "cp-all-in-one_default" with the default driver
Creating zookeeper ... done
Creating broker    ... done
Creating schema-registry ... done
Creating rest-proxy      ... done
Creating connect         ... done
Creating ksql-datagen    ... done
Creating ksql-server     ... done
Creating control-center  ... done
Creating ksql-cli        ... done
```

4. Optional: Run this command to verify that the services are up and running.

```
docker-compose ps
```

```
     Name                    Command               State                 Ports
-------------------------------------------------------------------------------
broker            /etc/confluent/docker/run        Up      0.0.0.0:29092->29092/tcp,
                                                           0.0.0.0:9092->9092/tcp
connect           /etc/confluent/docker/run        Up      0.0.0.0:8083->8083/tcp,
                                                           9092/tcp
control-center    /etc/confluent/docker/run        Up      0.0.0.0:9021->9021/tcp
ksql-cli          ksql http://localhost:8088       Up
ksql-datagen      bash -c echo Waiting for K ...   Up
ksql-server       /etc/confluent/docker/run        Up      0.0.0.0:8088->8088/tcp
rest-proxy        /etc/confluent/docker/run        Up      0.0.0.0:8082->8082/tcp
schema-registry   /etc/confluent/docker/run        Up      0.0.0.0:8081->8081/tcp
zookeeper         /etc/confluent/docker/run        Up      0.0.0.0:2181->2181/tcp,
                                                           2888/tcp, 3888/tcp
```

If the state is not `Up`, rerun the `docker-compose up -d` command.

Issue: Cannot

monitoring production data pipelines and event streaming applications.

1. Navigate to the Control Center web interface at http://localhost:9021/ (http://localhost:9021/) and select your cluster.

   Important

   It may take a minute or two for Control Center to come online.

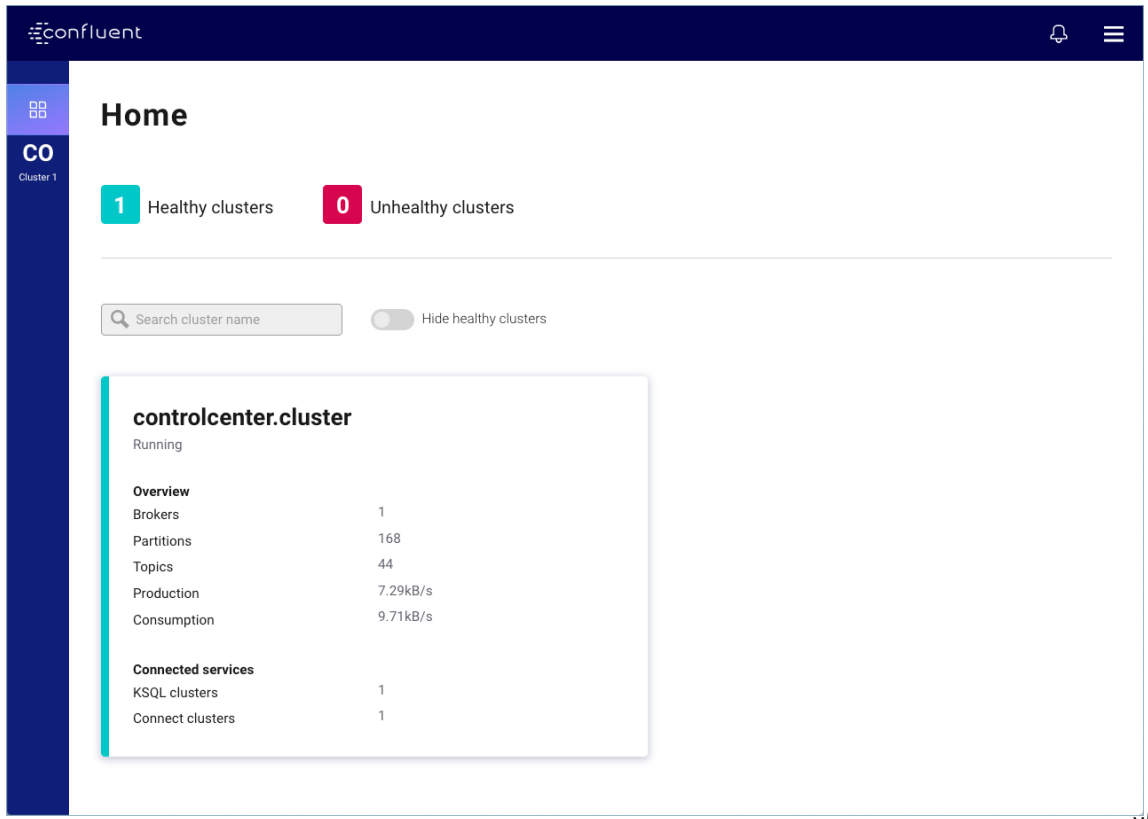

(../_images/c3-landing-page.png)

Issue: Cannot

2. Select **Topics** from the cluster submenu and click **Add a topic**.

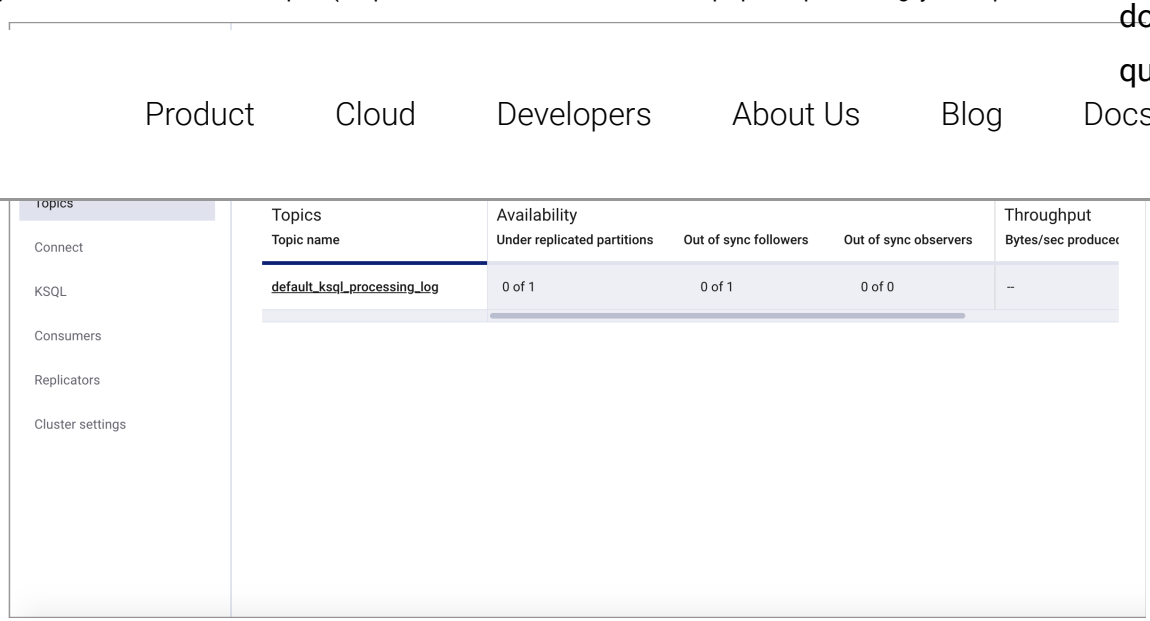

(../_images/c3-create-topic.png)

3. Create a topic named `pageviews` and click **Create with defaults**.

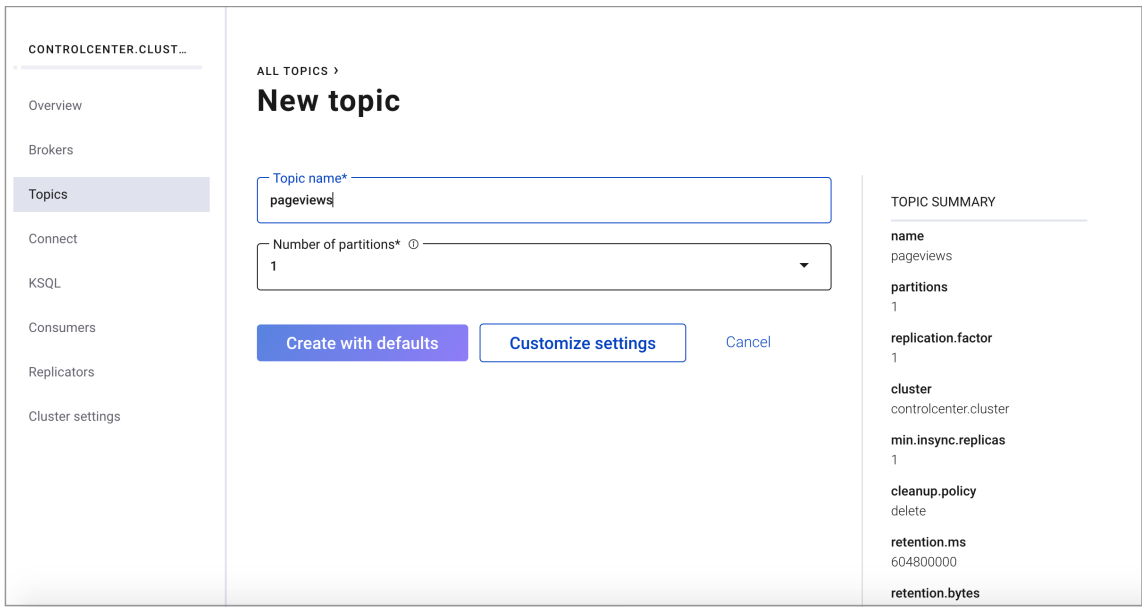

(../_images/c3-create-topic-name.png)

4. Repeat the previous steps and create a topic named `users` and click **Create with defaults**.

Issue: Cannot

Step 3: Install a Kafka Connector and Generate Sample Data

In this step, you use Kafka Connect to run a demo source connector called `kafka-connect-datagen` that creates sample data for the Kafka topics `pageviews` and `users`.

Tip

The Kafka Connect Datagen connector was installed automatically when you started Docker Compose with the `--build` argument in Step 1: Download and Start Confluent Platform Using Docker. If you encounter issues locating the Datagen Connector, refer to the Issue: Cannot locate the Datagen Connector in the Troubleshooting section.

1. Run one instance of the Kafka Connect Datagen (https://www.confluent.io/connector/kafka-connect-datagen/) connector to produce Kafka data to the `pageviews` topic in AVRO format.
    1. From your cluster, click **Connect**.
    2. Select the `connect-default` cluster and click **Add connector**.
    3. Find the DatagenConnector tile and click **Connect**.

        Tip

        To narrow displayed connectors, click **Filter by type** -> **Sources**.

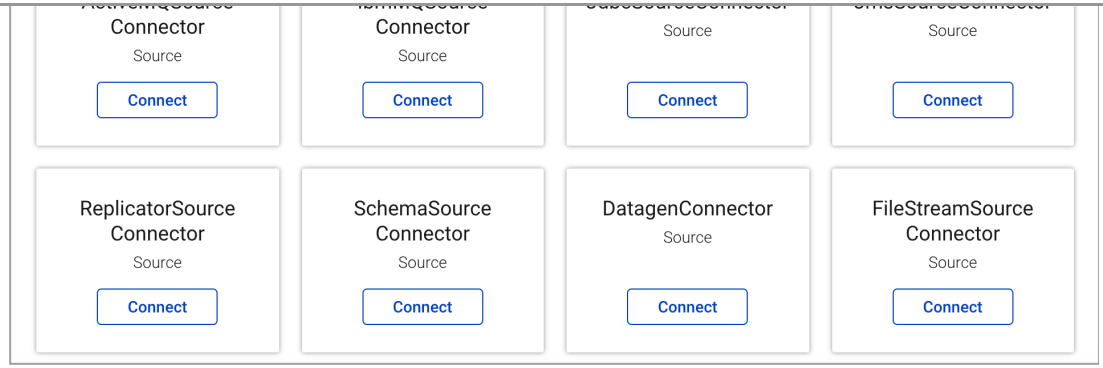

(../_images/connect-page-new-source.png)

4. Name the connector `datagen-pageviews`. After naming the connector, new fields appear. Scroll down and specify the following configuration values:

   - In the **Key converter class** field, type `org.apache.kafka.connect.storage.StringConverter`.
   - In the **kafka.topic** field, type `pageviews`.
   - In the **max.interval** field, type `100`.
   - In the **iterations** field, type `1000000000`.
   - In the **quickstart** field, type `pageviews`.

Issue: Cannot

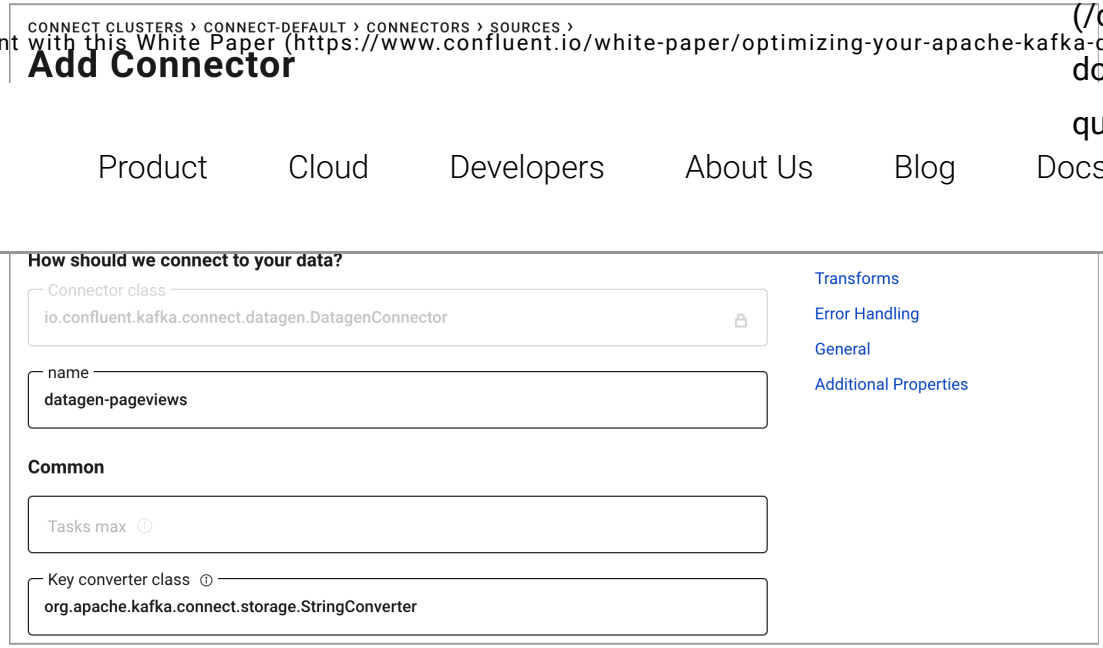

(../_images/connect-configure-pageviews.png)

5. Click **Continue**.
6. Review the connector configuration and click **Launch**.

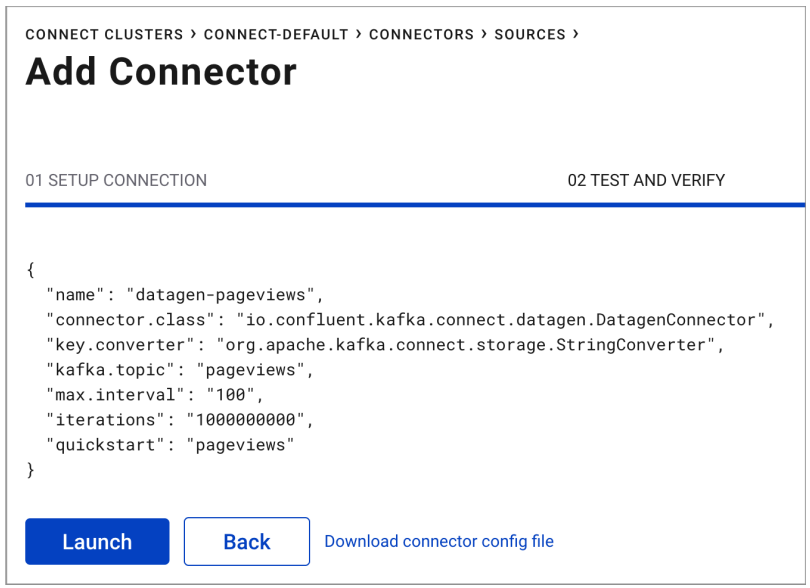

(../_images/connect-review-pageviews.png)

2. Run another instance of the Kafka Connect Datagen (https://www.confluent.io/connector/kafka-connect-datagen/) connector to produce Kafka

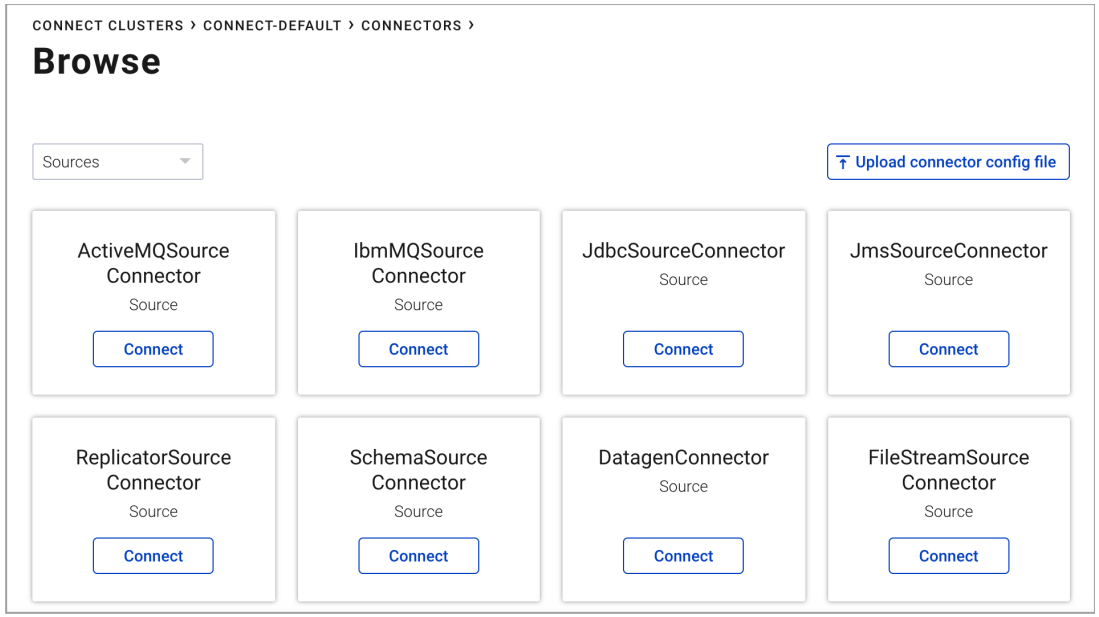

(../_images/connect-page-new-source.png)

3. Name the connector `datagen-users`. After naming the connector, new fields appear. Scroll down and specify the following configuration values:

   - In the **Key converter class** field, type `org.apache.kafka.connect.storage.StringConverter`.
   - In the **kafka.topic** field, type `users`.
   - In the **max.interval** field, type `1000`.
   - In the **iterations** field, type `1000000000`.
   - In the **quickstart** field, type `users`.

Issue: Cannot

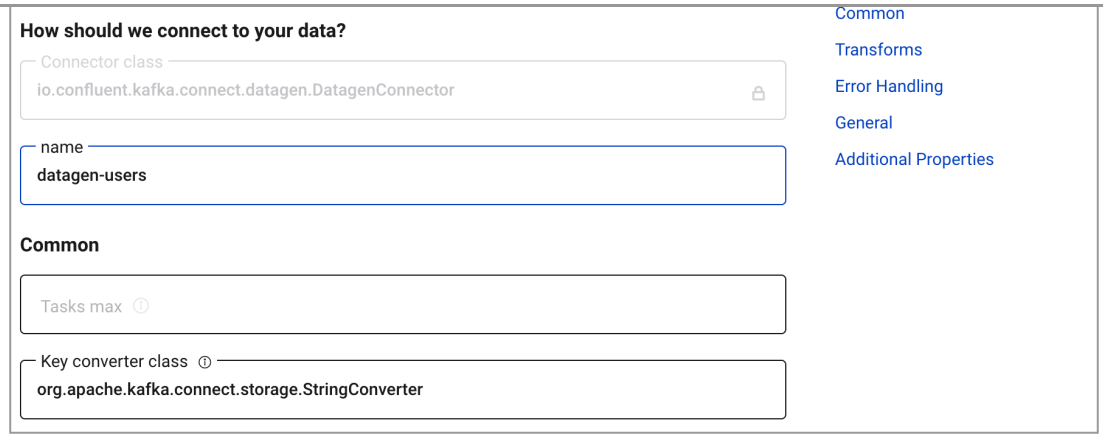

(../_images/connect-configure-users.png)

4. Click **Continue**.
5. Review the connector configuration and click **Launch**.

Issue: Cannot

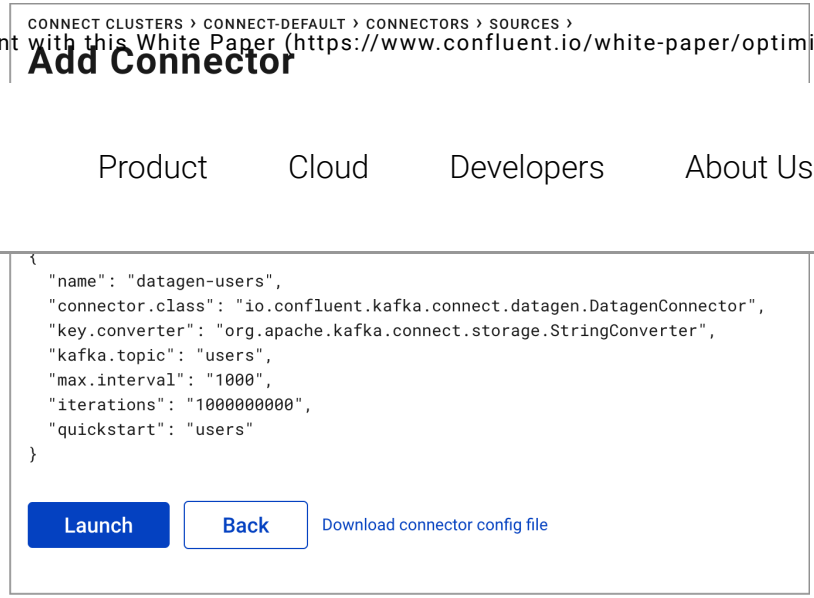

(../_images/connect-review-users.png)

Step 4: Create and Write to a Stream and Table using KSQL

In this step, KSQL queries are run on the `pageviews` and `users` topics that were created in the previous step. The KSQL commands are run using the KSQL tab in Control Center.

Tip

Issue: Cannot

You can also run these commands using the KSQL CLI (/ksql/docs/installation/index.html#install-ksql) from your Docker container with this (/current/quickstart/docker-docker-docker-quickstart.html#docker-quickstart)

Create Streams and Tables

In this step, KSQL is used to create a stream for the `pageviews` topic, and a table for the `users` topic.

1. From your cluster, click **KSQL** and choose the **KSQL** application.
2. From the **KSQL EDITOR** page, click the **Streams** tab and **Add Stream**.

(../_images/ksql-interface-create-stream2.png)

3. Select the `pageviews` topic.

Issue: Cannot

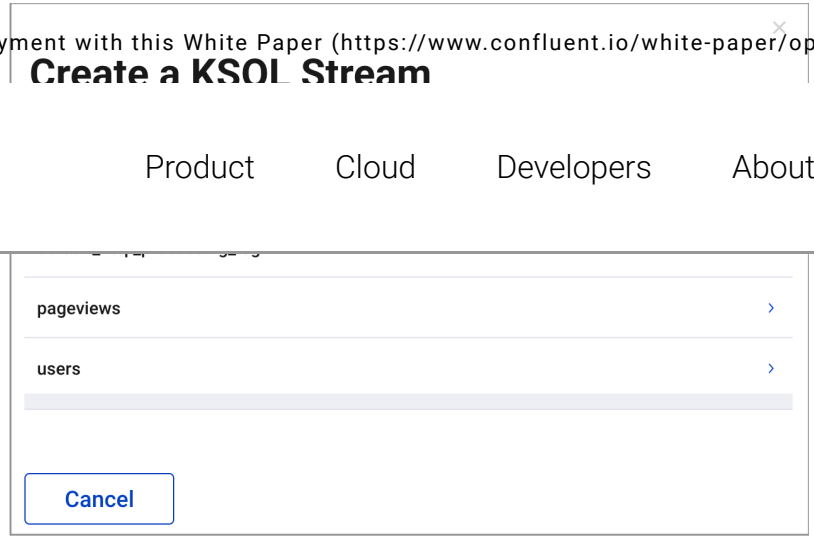

create-stream-pageview.png)

4. Choose your stream options:

   - In the **Encoding** field, select `AVRO`.
   - In the **Field(s) you'd like to include in your STREAM** field, ensure fields are set as follows:
     - `viewtime` with type `BIGINT`
     - `userid` with type `VARCHAR`
     - `pageid` with type `VARCHAR`

(../_images/c3-ksql-

Issue: Cannot

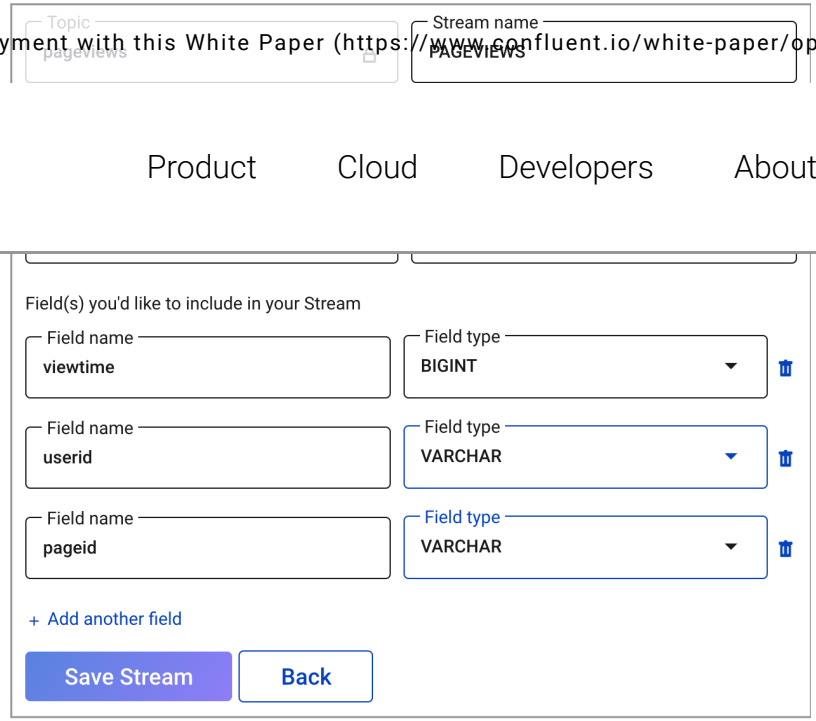

(../_images/c3-ksql-create-stream-pageview-2.png)

5. Click **Save Stream**.
6. Click the **Tables** tab -> **Add a Table** and select the `users` topic.

(../_images/c3-ksql-create-stream-users.png)

7. Choose your table options:

   - In the **Encoding** field, select `AVRO`.
   - In the **Key** field, select `userid`.
   - In the **Field(s) you'd like to include in your TABLE** field, ensure fields are set as follows:
     - `registertime` with type `BIGINT`
     - `userid` with type `VARCHAR`
     - `regionid` with type `VARCHAR`
     - `gender` with type `VARCHAR`

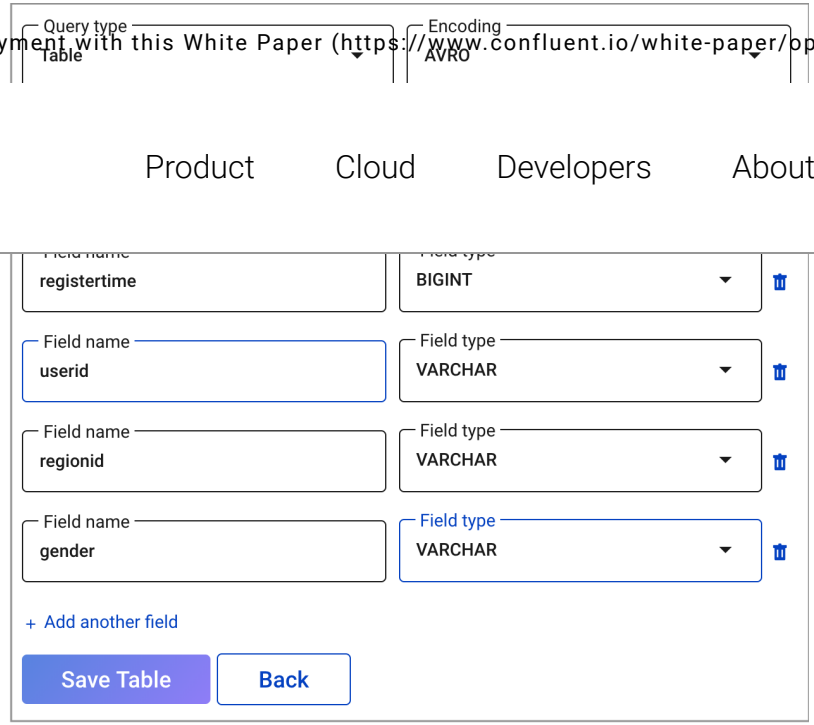

(../_images/c3-ksql-create-table-users.png)

8. Click **Save Table**.

Write Queries

These examples write queries using the **KSQL** tab in Control Center.

1. From your cluster, click **KSQL** and choose the **Editor** page.
2. From the **KSQL EDITOR** page, click **Add query properties** to add a custom query property. Set the `auto.offset.reset` parameter to `earliest`.

   This instructs KSQL queries to read all available topic data from the beginning. This configuration is used for each subsequent query. For more information, see the KSQL Configuration Parameter Reference (../ksql/docs/installation/server-config/config-

Issue: Cannot

reference.html#ksql-param-reference)

(/current/tutorials/docker-quickstart.html)

query-properties.png)

3. Run the following queries.

   1. Create a non-persistent query that returns data from a stream with the results limited to a maximum of three rows.

      ```
      SELECT pageid FROM pageviews EMIT CHANGES LIMIT 3;
      ```

      Your output should resemble:

      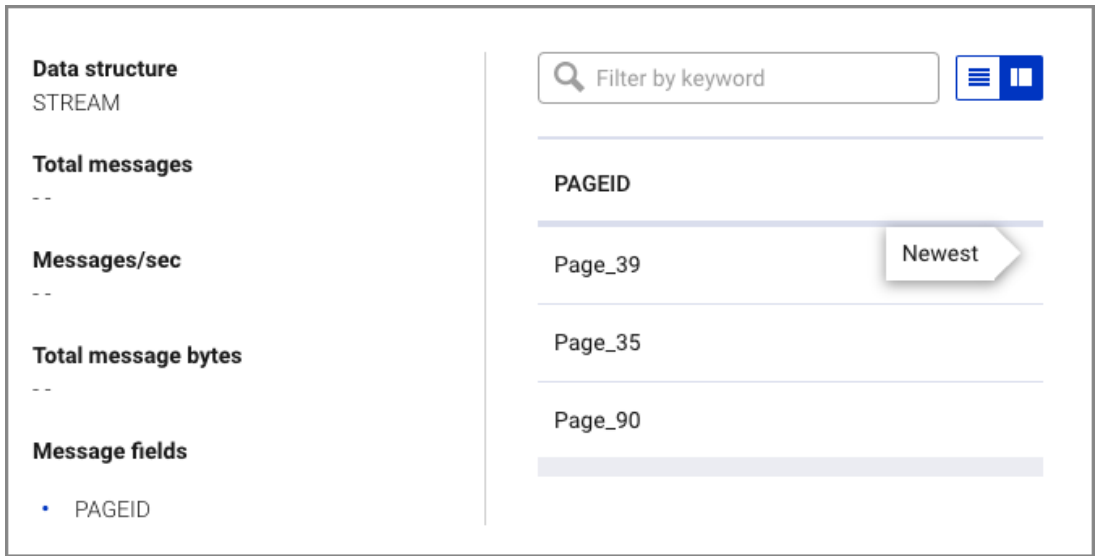

      (../_images/c3-ksql-query-results-pageid.png)

      Tip

Issue: Cannot

Click the **Card view** or **Tabular view** icon to change the layout. Click the expand icon to expand a message.

(../_images/c3-ksql-query-results-pageid-card.png)

2. Create a persistent query that filters for female users. The results from this query are written to the Kafka `PAGEVIEWS_FEMALE` topic. This query enriches the `pageviews` STREAM by doing a `LEFT JOIN` with the `users` TABLE on the user ID, where a condition (`gender = 'FEMALE'`) is met.

```
CREATE STREAM pageviews_female AS SELECT users.userid AS userid, pageid
, regionid, gender FROM pageviews LEFT JOIN users ON pageviews.userid =
users.userid WHERE gender = 'FEMALE';
```

Your output should resemble:

(../_images/c3-ksql-persist-query-pv-female-results.png)

3. Create a persistent query where a condition (`regionid`) is met, using `LIKE`. Results from this query are written to a Kafka topic named `pageviews_enriched_r8_r9`.

```
CREATE STREAM pageviews_female_like_89 WITH (kafka_topic='pageviews_enriched_r8_r9', value_format='AVRO') AS SELECT * FROM pageviews_female WHERE regionid LIKE '%_8' OR regionid LIKE '%_9';
```

Your output should resemble:

```
{
  "@type": "currentStatus",
  "statementText": "CREATE STREAM pageviews_female_like_89 WITH (kafka_topic='pageviews_enriched_r8_r9', value_format='AVRO') AS SELECT * FROM pageviews_female WHERE r
  "commandId": "stream/PAGEVIEWS_FEMALE_LIKE_89/create",
  "commandStatus": {
    "status": "SUCCESS",
    "message": "Stream created and running"
  }
}
```

(../_images/c3-ksql-persist-query-pv-female89-results.png)

4. Create a persistent query that counts the pageviews for each region and gender combination in a tumbling window (../streams/developer-guide/dsl-api.html#windowing-tumbling) of 30 seconds when the count is greater than 1. Because the procedure is grouping and counting, the result is now a table, rather than a stream. Results from this query are written to a Kafka topic called `PAGEVIEWS_REGIONS`.

```
CREATE TABLE pageviews_regions AS SELECT gender, regionid , COUNT(*) AS numusers FROM pageviews_female WINDOW TUMBLING (size 30 second) GROUP BY gender, regionid HAVING COUNT(*) > 1;
```


Issue: Cannot

Your output should resemble:

```
8 | )
```

(../_images/c3-ksql-persist-query-table-results.png)

5. Click **Running queries**. You should see the following persisted queries:

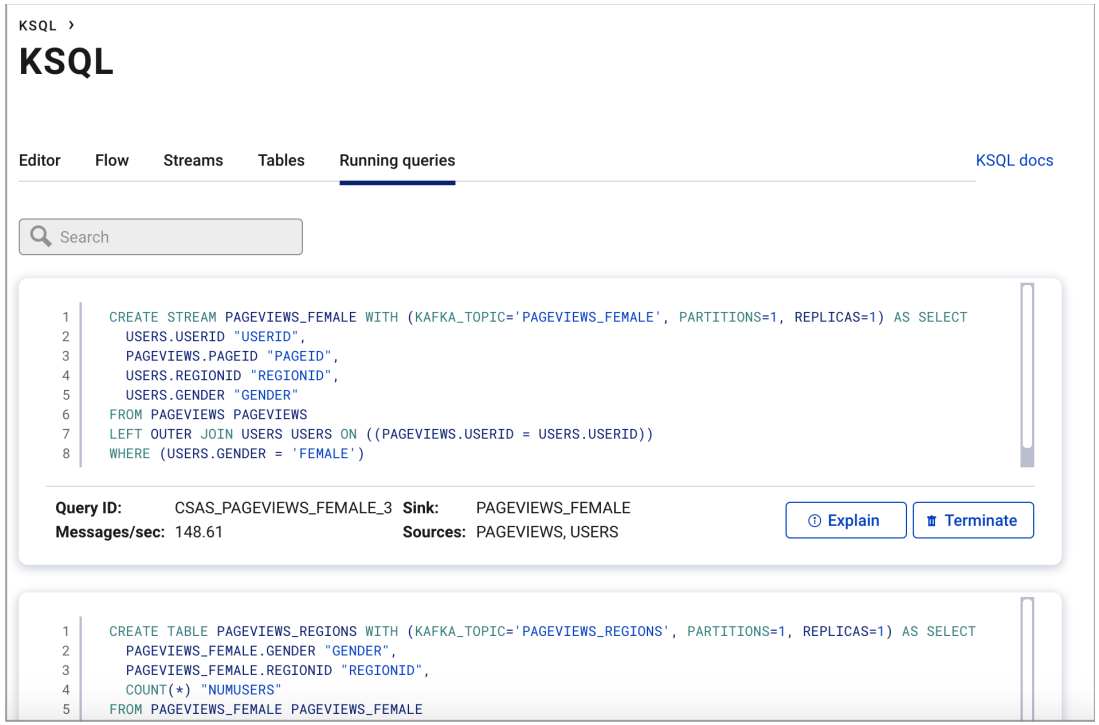

(../_images/c3-ksql-persistent-query1.png)

6. Click **Editor**. On the right side of the page, find the **All available streams and tables** pane, which shows all of the streams and tables that you can access.

Issue: Cannot

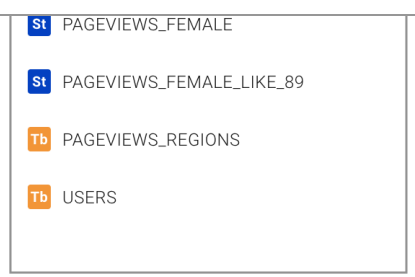

(../_images/c3-ksql-stream-table-view-1.png)

7. In the **All available streams and tables** section, click **KSQL_PROCESSING_LOG** to view the stream's schema, including nested data structures.

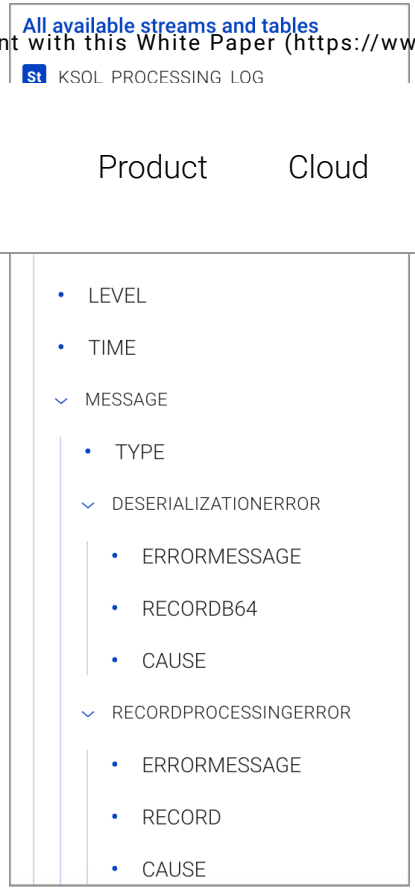

(../_images/c3-ksql-stream-table-view-2.png)

Step 5: Monitor Consumer Lag

Navigate to the **Consumers** tab to view the consumers created by KSQL.

Click the consumer group ID to view details for the `_confluent-ksql-default_query_CSAS_PAGEVIEWS_FEMALE_3` consumer group.

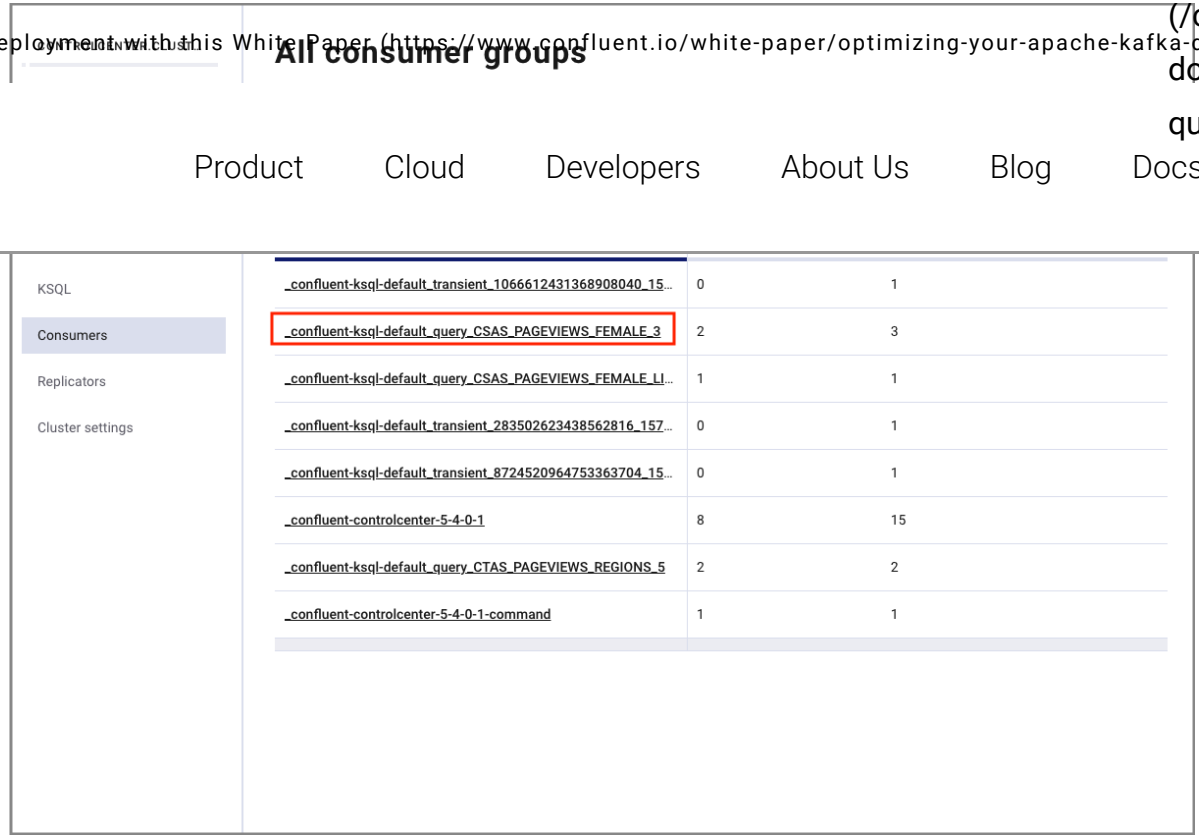

(../_images/ksql-interface-monitor.png)

From this page you can see the consumer lag and consumption values for your streaming query.

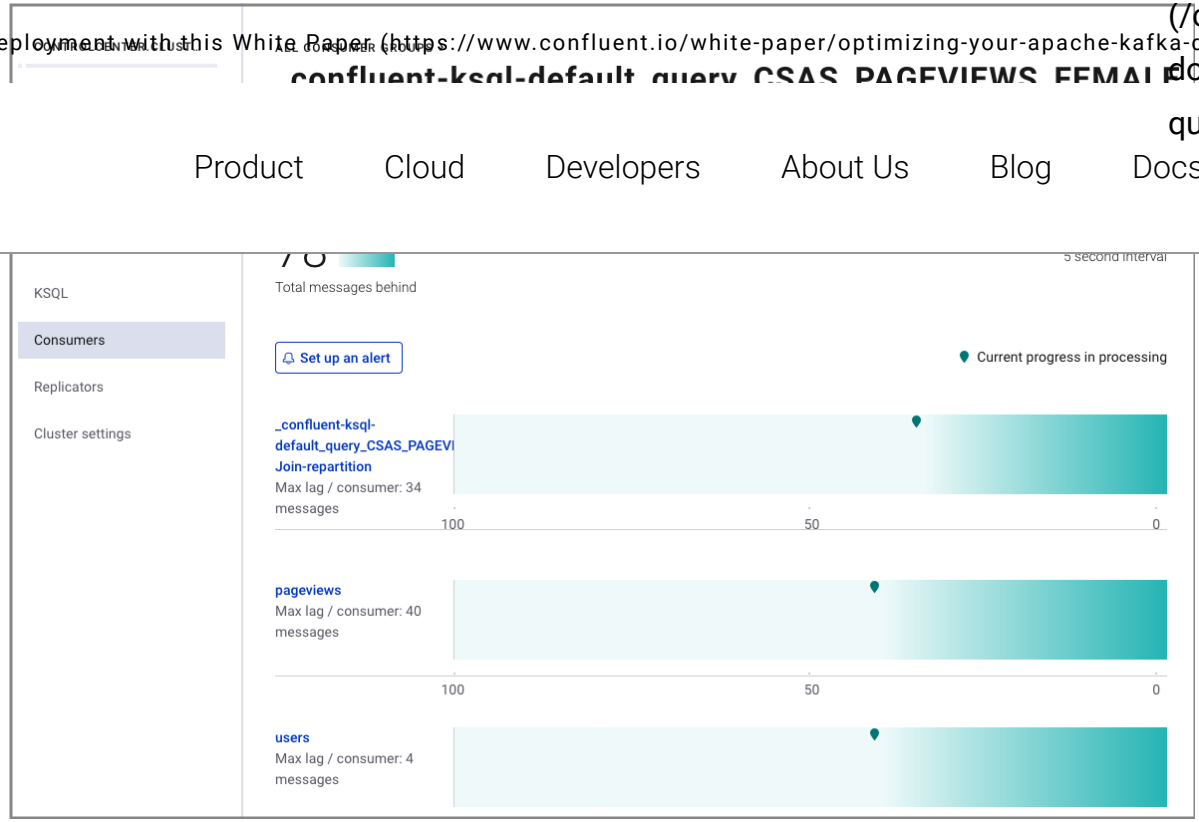

(../_images/ksql-interface-monitor-cnsmgp.png)

For more information, see the Control Center Consumers (../control-center/consumers.html#controlcenter-userguide-consumers) documentation.

Step 6: Stop Docker

When you are done working with Docker, you can stop and remove Docker containers and images.

Issue: Cannot

1. View a list of all Docker container IDs

2. Run the following command to stop the Docker containers for Confluent:

```
docker container stop $(docker container ls -a -q -f "label=io.confluent.docker")
```

3. Run the following commands to stop the containers and prune the Docker system. Running these commands deletes containers, networks, volumes, and images; freeing up disk space:

```
docker container stop $(docker container ls -a -q -f "label=io.confluent.docker") && docker system prune -a -f --volumes
```

Tip

Remove the filter label for Confluent Docker (`-f "label=io.confluent.docker"`) to clear all Docker containers from your system.

You can rebuild and restart the containers at any time using the `docker-compose up -d --build` command.

For more information, refer to the official Docker (https://docs.docker.com/) documentation.

Issue: Cannot

Issue: Cannot locate the Datagen Connector

Resolution: Make sure to run the `docker-compose` command with the `--build` option:

```
docker-compose up -d --build
```

For details, see Step 1: Download and Start Confluent Platform Using Docker.

Resolution: Run the `build` command just for connect.

```
docker-compose build --no-cache connect
```

Your output should resemble:

```
Building connect
...
Completed
Removing intermediate container cdb0af3550c8
---> 36d00047d29b
Successfully built 36d00047d29b
Successfully tagged confluentinc/kafka-connect-datagen:latest
```

Resolution: Check the Connect log for `Datagen`.

```
docker-compose logs connect | grep -i Datagen
```

```
connect  | [2019-04-17 20:03:26,137] INFO Loading plugin from: /usr/share/confluent-hub-components/confluentinc-kafka-connect-datagen (org.apache.kafka.connect.runtime.isolation.DelegatingClassLoader)
connect  | [2019-04-17 20:03:26,206] INFO Registered loader: PluginClassLoader{pluginLocation=file:/usr/share/confluent-hub-components/confluentinc-kafka-connect-datagen/} (org.apache.kafka.connect.runtime.isolation.DelegatingClassLoader)
connect  | [2019-04-17 20:03:26,206] INFO Added plugin 'io.confluent.kafka.connect.datagen.DatagenConnector' (org.apache.kafka.connect.runtime.isolation.DelegatingClassLoader)
connect  | [2019-04-17 20:03:28,102] INFO Added aliases 'DatagenConnector' and 'Datagen' to plugin 'io.confluent.kafka.connect.datagen.DatagenConnector' (org.apache.kafka.connect.runtime.isolation.DelegatingClassLoader)
```

Resolution: Check the Connect log for a warning and reminder to run the `docker-compose up -d --build` command properly.

```
docker-compose logs connect | grep -i Datagen
```

If the following warning is present, re-run the `docker-compose up -d --build` command:

```
connect  | WARNING: Did not find directory for kafka-connect-datagen (did you remember to run: docker-compose up -d --build ?)
```

Resolution: Verify the `.jar` file for `kafka-connect-datagen` has been added and is present in the `lib` subfolder.

```
docker-compose exec connect ls /usr/share/confluent-hub-components/confluentinc-kafka-connect-datagen/lib/
```

Resolution: Verify the plugin exists in the connector path.

```
docker-compose exec connect bash -c 'echo $CONNECT_PLUGIN_PATH'
```

Your output should resemble:

```
/usr/share/java,/usr/share/confluent-hub-components
```

Confirm its contents are present:

```
docker-compose exec connect ls /usr/share/confluent-hub-components/confluentinc-kafka-connect-datagen
```

Your output should resemble:

```
assets   doc  etc  lib  manifest.json
```

Resolution: In **Kafka Connect > Setup Connection**, scroll down through the list of connectors to locate **DatagenConnector**; there are multiple connectors in the menu.

Issue: Stream-Stream joins error

An error states Stream-Stream joins must have a WITHIN clause specified. This error can occur if you created streams for both `pageviews` and `users` by mistake.

(../_images/c3-ksql-stream-stream-join-error.png)

Resolution: Ensure that you created a *stream* for `pageviews`, and a *table* for `users` in Step 4: Create and Write to a Stream and Table using KSQL.

Issue: Unable to successfully complete KSQL query steps

Java errors or other severe errors were encountered.

Resolution: Ensure you are on an Operating System (../installation/versions-interoperability.html#operating-systems) currently supported by Confluent Platform.

Resolution: Ensure that the Docker memory was increased to 8 MB. Go to **Docker > Preferences > Advanced**. If Docker memory is insufficient, other unpredictable issues could occur.

Next Steps

Learn more about the components shown in this quick start:

- KSQL documentation (../ksql/docs/index.html#ksql-home) Learn about processing your data with KSQL for use cases such as streaming ETL, real-time monitoring, and anomaly detection. You can also learn how to use KSQL with this collection of scripted demos (https://github.com/confluentinc/examples).
- Stream Processing Cookbook (https://www.confluent.io/stream-processing-cookbook/) Try out in-depth KSQL tutorials and recommended deployment scenarios.
- Kafka Streams documentation (../streams/index.html#kafka-streams) Learn how to build stream processing applications in Java or Scala.
- Kafka Connect documentation (../connect/index.html#kafka-connect) Learn how to integrate Kafka with other systems and download ready-to-use connectors (https://www.confluent.io/product/connectors/) to easily ingest data in and out of Kafka in real-time.
- Kafka Clients documentation (../clients/index.html#kafka-clients) Learn how to read and write data to and from Kafka using programming languages such as Go, Python, .NET, C/C++.
- Videos, Demos, and Reading Material (../tutorials/index.html#tutorials) Try out the Confluent Platform tutorials and examples, watch demos and screencasts, and learn with white papers and blogs.

Last updated on Feb 26, 2020.